CS 553 - Cloud Computing, Spring 2018

PA 1 - Read Me Manual

Benchmarking

Saumil Pareshbhai Ajmera | A20397303 | sajmera4@hawk.iit.edu

This document will help in running each application on cluster and helps to know how input is passed and where output is obtained. Also it includes screenshots of benchmarks ran on cluster.

1. Processor Benchmarking

Steps to run MyCPUBench: Hyperion Cluster

1) Run the *Makefile* which will compile the *MyCPUBench.c* code with required flags into executable *MyCPUBench*
2) Run script file *create_slurms* using *bash create_slurms* command which will generate config files naming like *cpu1.slurm*, *cpu2.slurm* etc. which contains slurm commands to run exe by passing required input arguments as input DAT file and output DAT file.
3) Now I have prepared single script file per 9 config files to submit 9 jobs to different nodes. Run1.sh and Run.sh are two script files.
4) Now do *sbatch run1.sh* and after 9 files has been processed do *sbatch run2.sh* for running remaining jobs.
5) Output is written into file into output folder named *cpu_SP_1thread.out.dat* with required set of values.

Steps to run Linpack Benchmark:

1) Go to the path mentioned below in the linpack folder after unzipping the tar file - *Linpack /l_mklb_p_2018.0.006/benchmarks_2018/linux/mkl/benchmarks/linpack*
2) Modify the *lininput_xeon64* file by number of trials, threads, problem size and ¾ of pmemory.

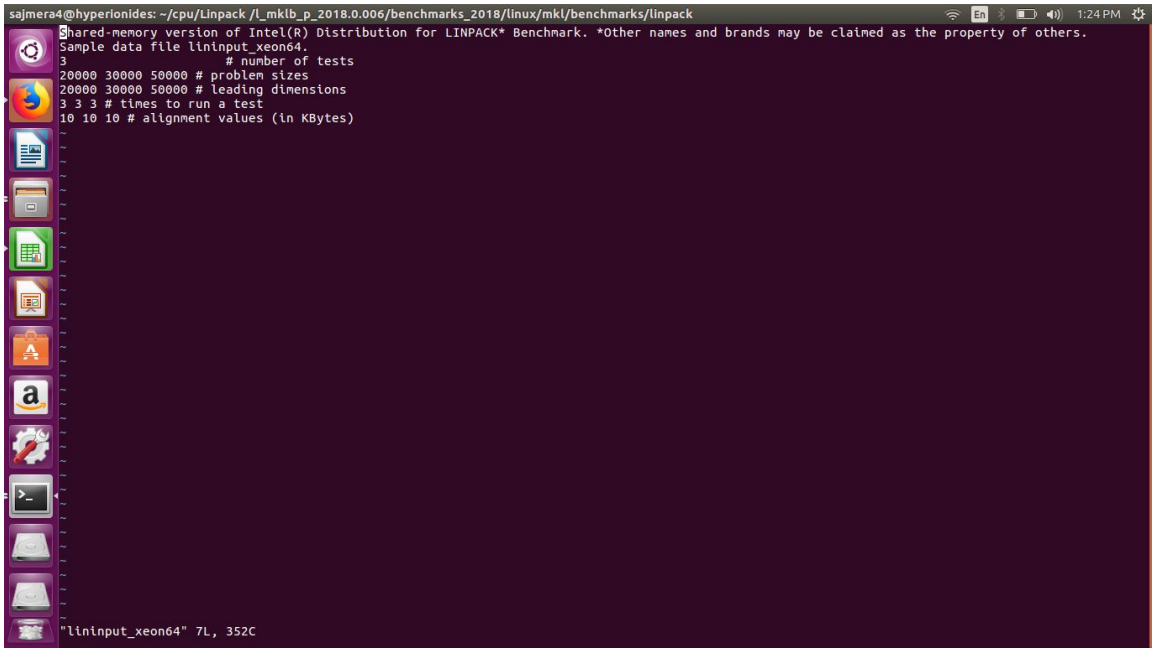

3) Now create run1.sh file as per below screenshot to call executable runme_xeon64

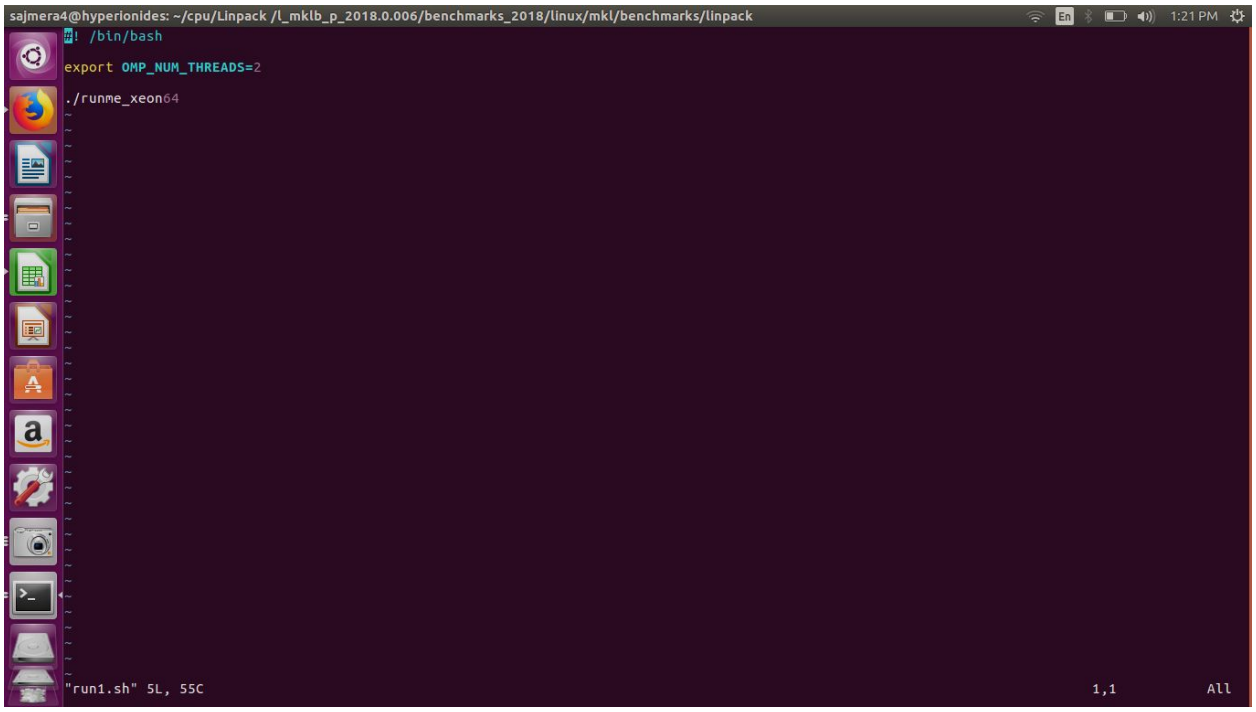

4) Output is written into file into slurm.out file in the same folder with following values of Giga Ops.

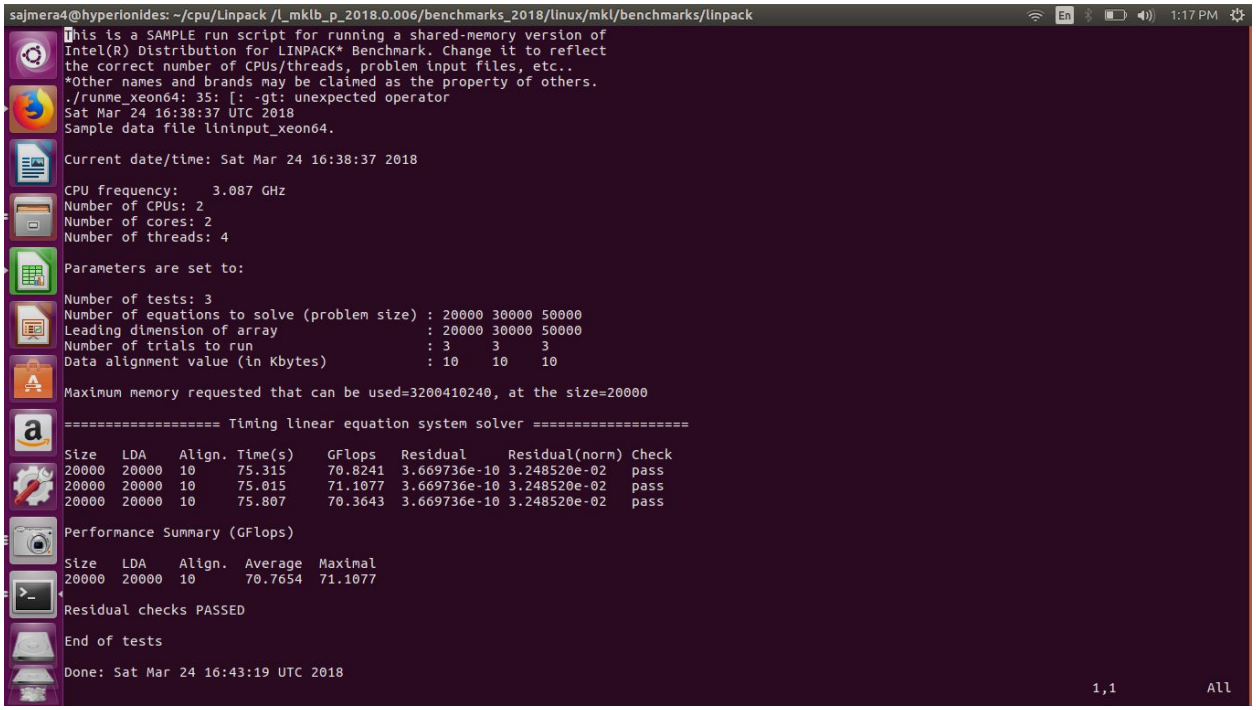

2. Memory Benchmarking

Steps to run MyRAMBench: - Hyperion Cluster

1) Run the *Makefile* which will compile the *MyRAMBench.c* code with required flags into executable *MyRAMBench*
2) Run script file *create_slurms* using *bash create_slurms* command which will generate config files naming like *mem1.slurm*, *mem2.slurm* etc. which contains slurm commands to run exe by passing required input arguments as input DAT file and output DAT file.
3) Now I have prepared single script file per 9 config files to submit 9 jobs to different nodes. Run1.sh, Run2.sh and Run3.sh script files.
4) Now do *sbatch run1.sh* and after 9 files has been processed do *sbatch run2.sh* for running remaining jobs. Similarly for other script files
5) Output is written into file in the output folder named *memory-RWR-1-1thread.out.dat* with required set of values. Output for latency is written into file in the output folder name *memory-latency.out.dat* with required set of values.

Steps to run PMBW Benchmark:

1) I have made one file pmbw1.sh which contains command to call executable of pmbw by passing certain parameters for our criteria.

2) -p and -P are min and max threads, -s and -S are min and max block size , -M for main block of memory and *PermRead64SimpleLoop* and *ScanRead64PtrSimpleLoop* are used for RWR and RWR throughput calculation. Here bandwidth term in generated output file gives the throughput value.

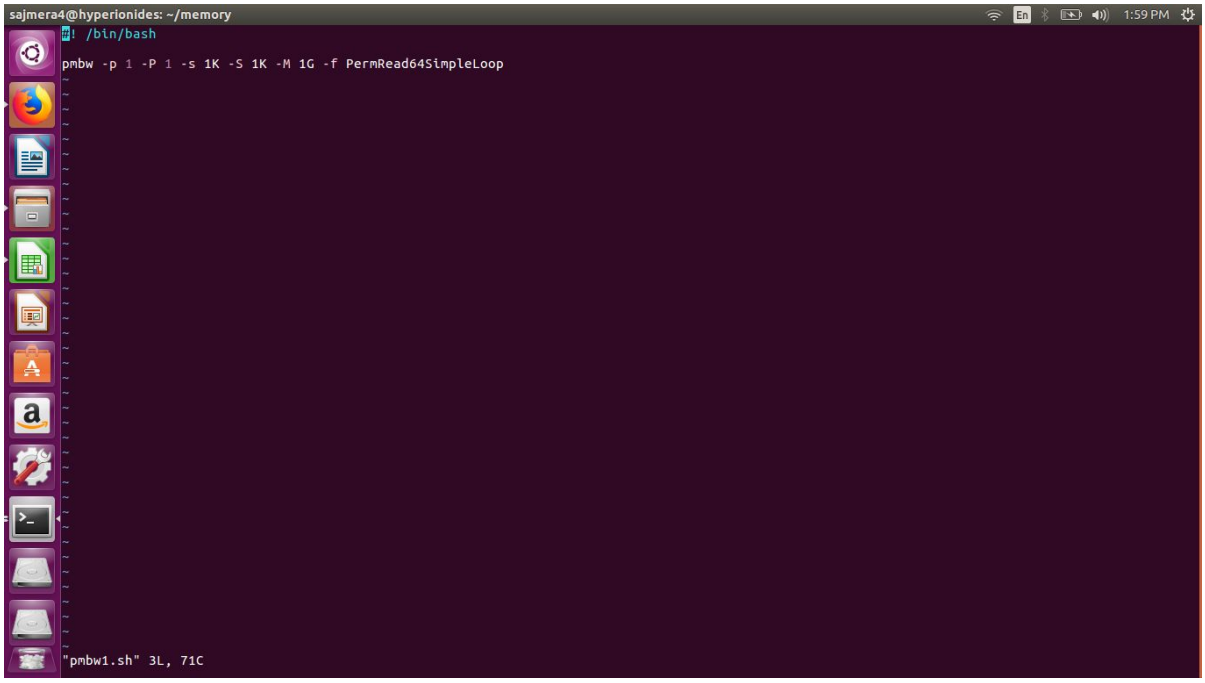

3) For latency calculation *ScanRead32PtrSimpleLoop* is used to find latency for 1B of data, its value is stored in rate format in the generated output file.

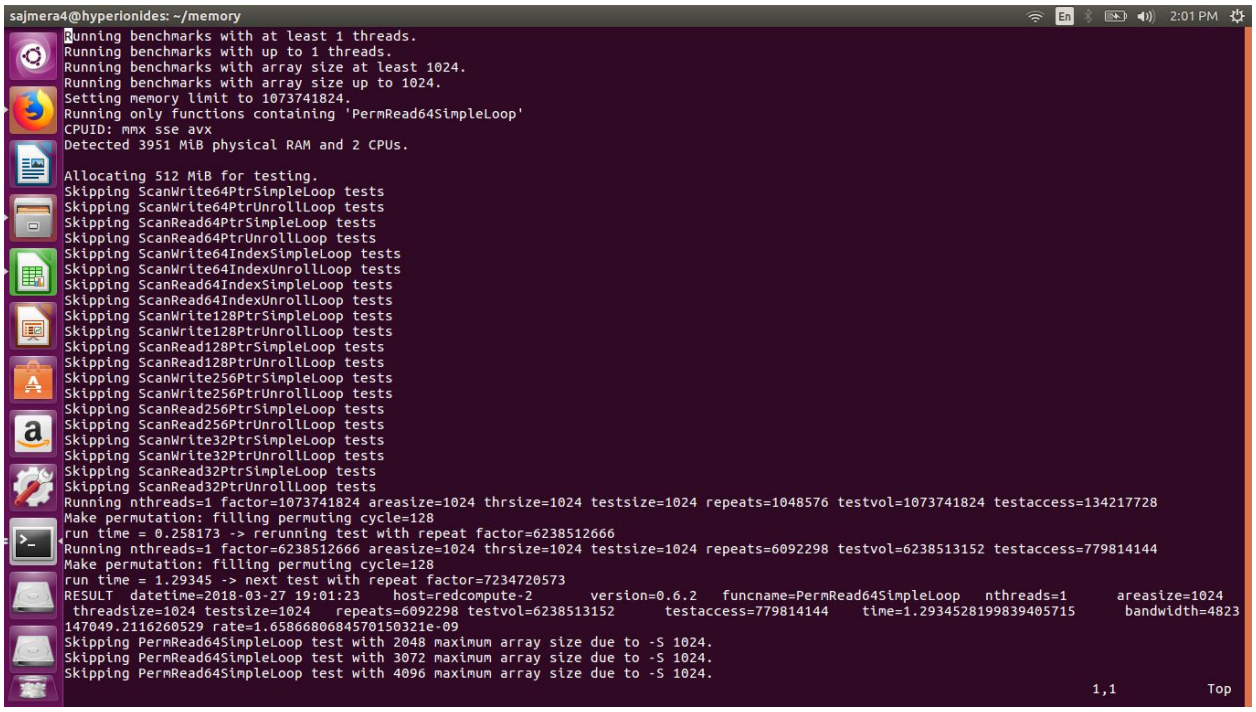

3. Disk Benchmarking

Steps to run MyDiskBench: Prometheous Cluster

1) Run the *Makefile* which will compile the *MyDiskBench.c* code with required flags into executable *MyDiskBench*
2) Run script file *create_slurms* using *bash create_slurms* command which will generate config files naming like *disk1.slurm*, *disk2.slurm* etc. which contains slurm commands to run exe by passing required input arguments as input DAT file and output DAT file.
3) Now I have prepared single script file per 9 config files to submit 9 jobs to different nodes. Here there are 52 experiments so there are 9 run.sh files containing all jobs.
4) Now do *sbatch run1.sh* and after 9 files has been processed do *sbatch run2.sh* for running remaining jobs. Similarly for other script files
5) Output is written into file in the output folder named *memory-RR-1-1thread.out.dat* with required set of values. Output for latency is written into file in the output folder name *disk-latency-IOPS.out.dat* with required set of values.

Steps to run IOZone Benchmark:

1) I have made one file runIozone.sh which contains command to call executable of iozone by passing certain parameters for our criteria.
2) -s indicates file size, -r indicates block size, -t indicates throughput in parallel threads -i indicates operation type 0,1,2 for sequence read, sequence write and random read write, -F indicates file name to be generated.

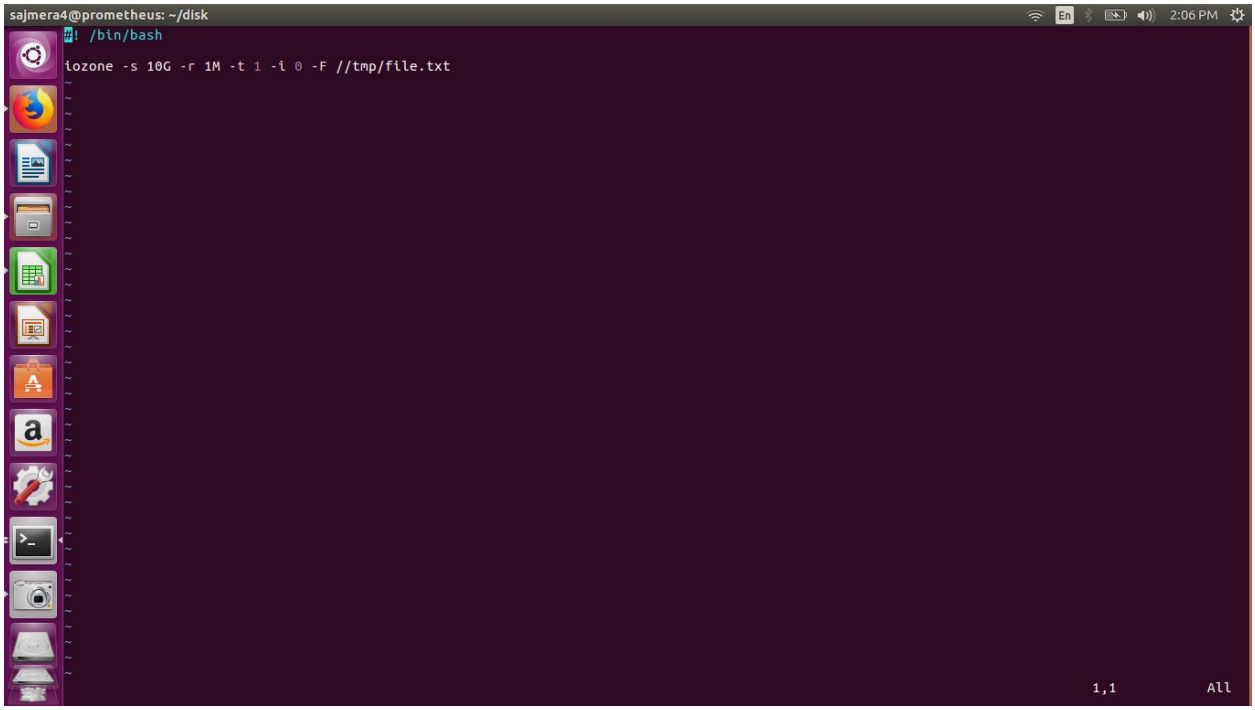

3) Here output file is generated after running sbatch runIozone.sh as slurm.out with following details.

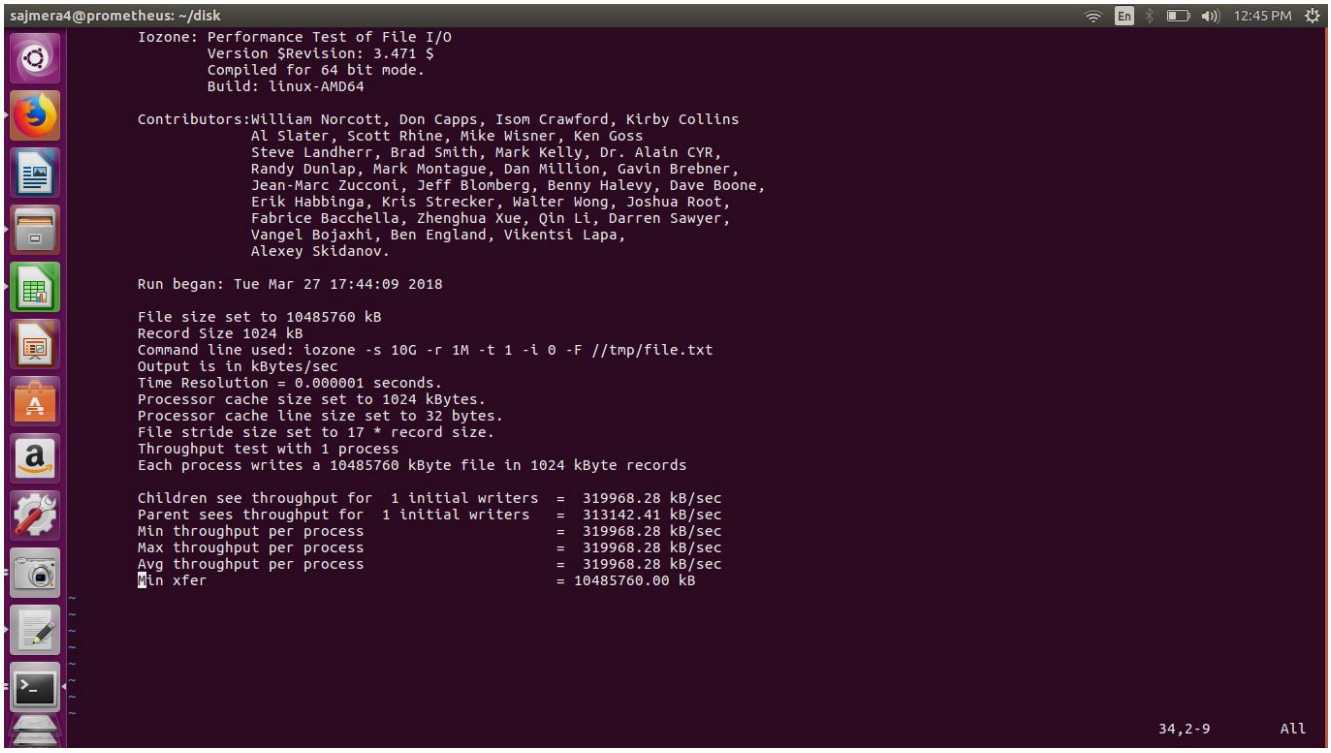

4. Network Benchmarking

Steps to run MyNETBench TCP/UDP : Hyperion Cluster

1) Run the *Makefile* which will compile the *MyNETBench-TCP.c and My NETBench-UDP.c* code with required flags into executable *MyNETBench-TCP and MyNETBench-UDP.*
2) Run script file *create_slurms* using *bash create_slurms* command which will generate config files for each configurations. which contains slurm commands to srun runUDP.sh or runTCP.sh by passing required input arguments as input DAT file and output DAT file.
3) Now manually one needs to submit each of file like sbatch TCP-1-1.slurm which will call runTCP.sh or (runUDP.sh) which will run same executable one for server and one for client by passing S or C for client server identification.
4) Note that in runTCP.sh and runUDP.sh , I have specifically mentioned path **//exports/home/sajmera4/network/MyNETBench-TCP,** **//exports/home/sajmera4/network/MyNETBench-UDP,** This may require changes depending upon how you are running.
5) Output for TCP and UDP for throughput is written into file in the output folder named *network-TCP-1thread.out.dat* with required set of values. Output for latency IOPS is written into file in the output folder name *network-latency-.out.dat* with required set of values.

Steps to run Iperf Benchmark:

1) I have made one file *iperf.slurm* which contains command to call *iperf.sh* twice for server and client.
2) Iperf.sh contains script to extract compute node and pass as parameter to client for running client and server with -P as threads, -s and -c for Server and Client , -u for UDP, -l for block size.

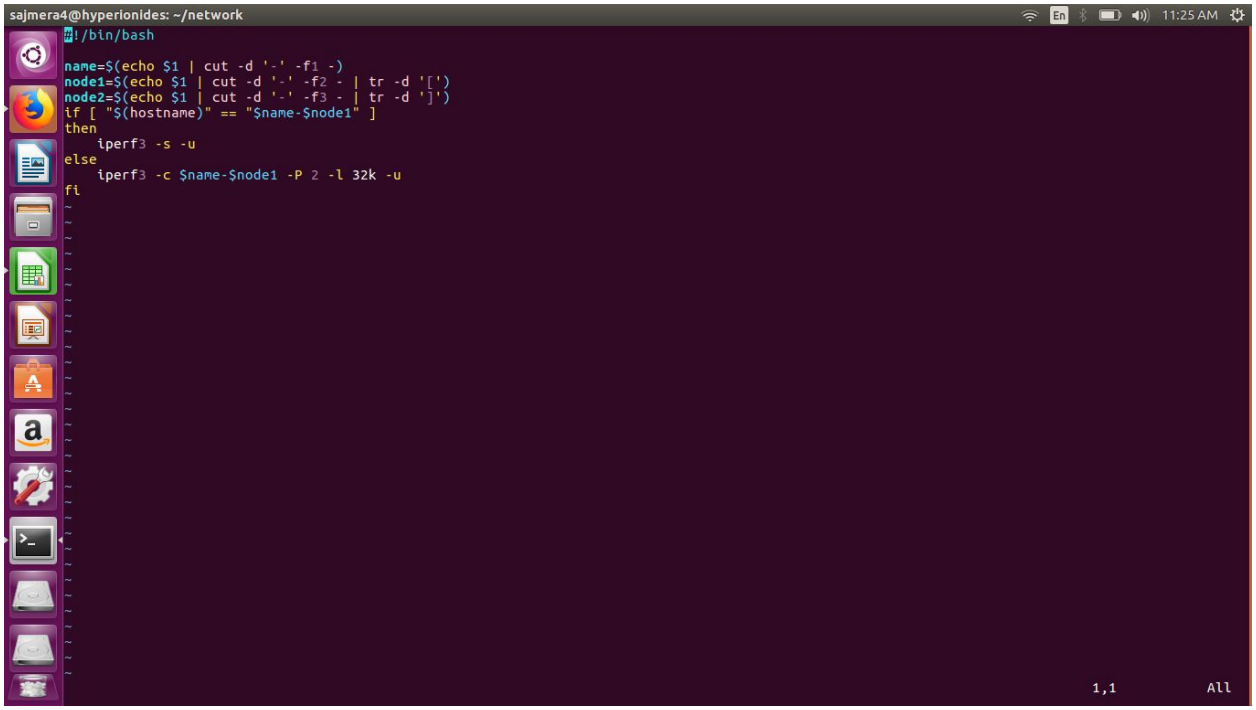

3) Here below screenshot shows the output of the script and it is stored in *iperf.out.*

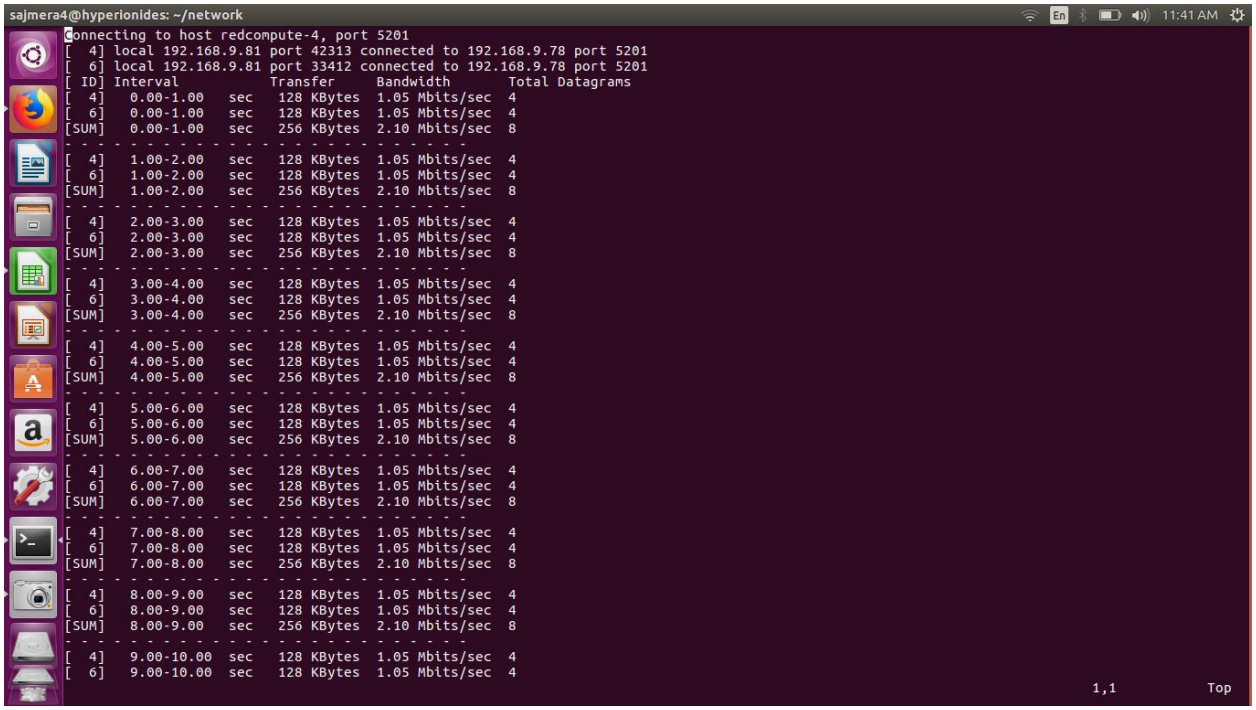

4) Ping Utility script - I have used ping command to test 1 byte of data from client to server and to run it do *sbatch ping.slurm* which will call *ping.sh* twice where different conditions are set. -s for buffer size -w for timeout .

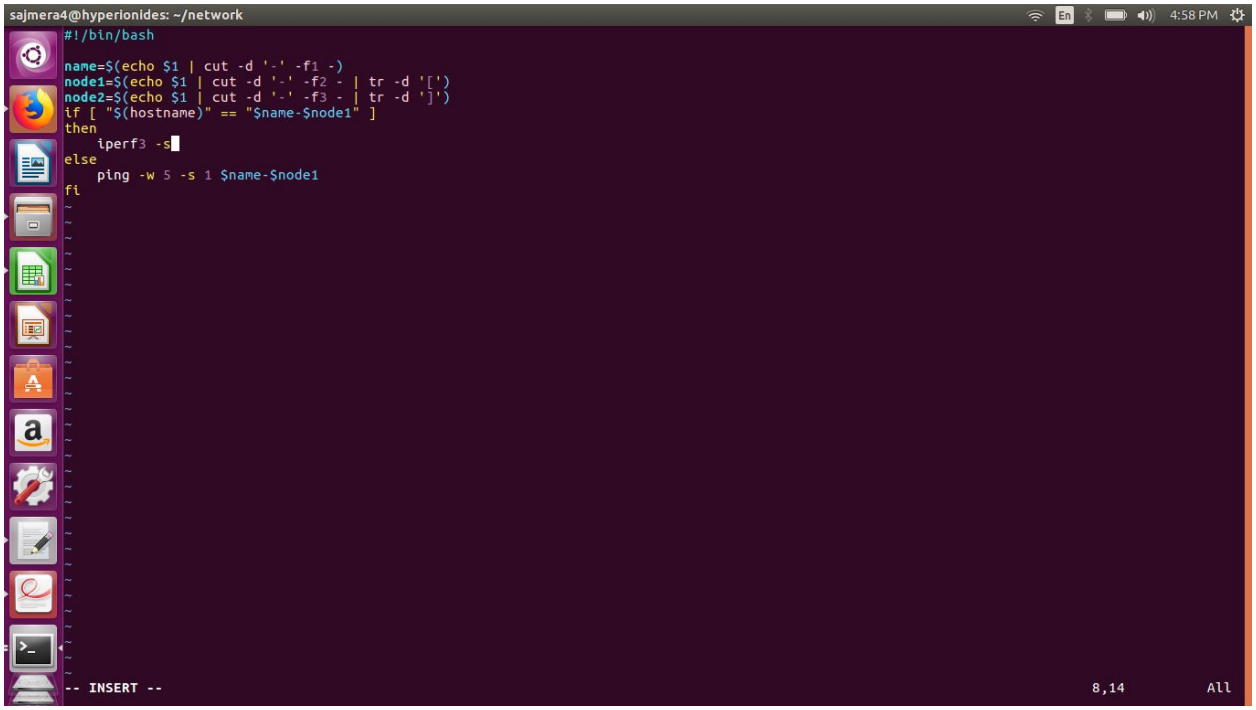

5) Output is stored in *ping.out*